

SQL: Null Values

Missing information

In a real database, we may have some missing information. For example, suppose we had a `spouse` column in the `Student` table. What should we put in this column for students who are not married? This isn't actually a case of missing information, it is a case of bad schema design! Since many or most students are not married, it doesn't make sense to have a column for spouses in the `Student` table. Instead, we can put it in another table with this schema (boldface is used to indicate that attributes `student1` and `student2` together form a key):

Married(student1, student2)

Only students who actually are married would appear in this table.

Sometimes we really do have missing information. For example, not every student has a spouse, but every student does belong to a specific campus. But suppose we have someone in the `Student` table whose campus we don't know. There is a correct campus value for this student, we just don't know it. What value should we put there?

In SQL, we use a special value called `NULL`.

Working with NULL

We can get a `NULL` value into a table with `INSERT INTO`. Recall the `Student` table:

```
csc343h-prof=> SELECT *
csc343h-prof-> FROM Student;
  sid | firstname | surname | campus | email | cgpa
-----+-----+-----+-----+-----+-----
 99132 | Avery     | Marchmount | StG    | avery@cs | 3.13
 98000 | William  | Fairgrieve | StG    | will@cs | 3.90
 99999 | Afsaneh  | Ali       | UTSC   | aali@cs | 2.98
 157  | Leilani  | Lakemeyer | UTM    | lani@cs | 3.42
 11111 | Homer    | Simpson   | StG    | doh@gmail | 0.40
(5 rows)
```

Here we insert a student with no value for first name or campus:

```
INSERT INTO Student VALUES (35461, NULL, 'Batman', NULL, 'bruce@batcave.com', 4.0);
```

If we don't want to allow `NULL` values in a given column, we can prevent them by using a `NOT NULL` constraint in the table definition. (We'll learn about table definitions a little later.) In fact, our schema was written so that `NULL` is not permitted in either column `firstname` or `campus`, so this insertion is unsuccessful.

We can compare a value to `NULL` using `IS NULL`. In our schema, the `breadth` column in table `Course` allows `NULL` values, so this query may yield some results:

```
csc343h-prof=> SELECT *
csc343h-prof-> FROM Course
csc343h-prof-> WHERE breadth IS NULL;
 cnum |          name          | dept | breadth
-----+-----+-----+-----
  100 | CSC for future prime ministers | CSC  |
(1 row)
```

Notice that the `NULL` value in column `breadth` is presented in the output as whitespace.

We can also compare a value to `NULL` using `IS NOT NULL`:

```
csc343h-prof=> SELECT *
csc343h-prof-> FROM Course
csc343h-prof-> WHERE breadth IS NOT NULL;
 cnum |          name          | dept | breadth
-----+-----+-----+-----
  343 | Intro to Databases      | CSC  | f
  207 | Software Design         | CSC  | f
  148 | Intro to Comp Sci       | CSC  | f
  263 | Data Struct & Anal      | CSC  | f
  320 | Intro to Visual Computing | CSC  | f
  200 | Intro Archaeology       | ANT  | t
  203 | Human Biol & Evol       | ANT  | f
  150 | Organisms in Environ    | EEB  | f
  216 | Marine Mammal Bio       | EEB  | f
  263 | Compar Vert Anatomy     | EEB  | f
  110 | Narrative               | ENG  | t
  205 | Rhetoric                | ENG  | t
  235 | The Graphic Novel       | ENG  | t
  200 | Environmental Change    | ENV  | f
  320 | Natl & Intl Env Policy  | ENV  | f
  220 | Mediaeval Society       | HIS  | t
  296 | Black Freedom           | HIS  | t
  222 | COBOL programming       | CSC  | f
(18 rows)
```

NULL introduces a third truth value: unknown

Because of the possible presence of **NULL** values, sometimes we don't know if a condition is true or false! For example, consider:

```
SELECT *
FROM Student
WHERE campus = 'StG';
```

If we have a row with a **NULL** value for **campus**, we cannot determine whether or not it is **StG**. In SQL, **if one or both operands is NULL, a comparison evaluates to "unknown"**.

Let's look at the implications of having this third truth-value.

Impact of NULL values on WHERE

WHERE is picky: it will not include a row for which the truth-value of the **WHERE** clause is unknown. So in our example above, if we didn't know a student's campus, they would not be included in the result of the query.

NATURAL JOIN is picky also. You can remember this if you think of a natural join as being a Cartesian product followed by a **WHERE** condition.

Impact of NULL values on aggregation

When we aggregate on a column, for example to find the average, **NULL** values in that column are ignored:

- NULL** never contributes to a **sum**, **avg**, or **count**
- NULL** makes no difference to the **max** or **min** of a column, unless every value is **NULL**. In that case, both the **max** and the **min** are **NULL**.

Here's a summary:

	some nulls in A	All nulls in A
min(A)	ignore the nulls	null
max(A)		
sum(A)		
avg(A)		
count(A)	all tuples count	0
count(*)		

Let's look at an example that uses this table about students at Runnymede Public School:

```
csc343h-prof=> -- Column grade has one NULL,
```

```
csc343h-prof=> -- and column age is entirely NULLS.
csc343h-prof=> SELECT *
csc343h-prof-> FROM Runnymede;
  name | age | grade
-----+-----+-----
diane |    |      8
will  |    |      8
cate  |    |      1
tom   |    |
micah |    |      1
grace |    |      2
(6 rows)
```

Tom's **NULL** grade doesn't contribute to any of the aggregations below. But **count(*)** counts rows, and it includes every row, regardless of any **NULL** s.

```
csc343h-prof=> SELECT min(grade), max(grade), sum(grade), avg(grade),
csc343h-prof=>          count(grade), count(*)
csc343h-prof-> FROM Runnymede;

 min | max | sum |          avg          | count | count
-----+-----+-----+-----+-----+-----
   1 |   8 |  20 | 4.0000000000000000 |     5 |     6
(1 row)
```

If an *entire* column is **NULL**, it can't have a meaningful **avg**, for example. The value of any aggregation on that column is **NULL**. There is one exception: the **COUNT** of such a column is **0**. Let's aggregate on our entirely **NULL** column, **age**.

```
csc343h-prof=> SELECT min(age), max(age), sum(age), avg(age),
csc343h-prof=>          count(age), count(*)
csc343h-prof-> FROM Runnymede;

 min | max | sum | avg | count | count
-----+-----+-----+-----+-----+-----
    |    |    |    |     0 |     6
(1 row)
```

Notice that **count(*)** gave the same answer as before. and that **count(age)** gave 0 because no values in that column are non-null.

There are corner cases that we haven't examined. For example, are two **NULL** values considered distinct from each other when we do **NATURAL JOIN**, set operations, or **SELECT DISTINCT**? This behaviour may vary across DBMSs, so it is best not to assume you know what will happen. Check the documentation.

Impact of NULL values on boolean conditions

Now that we have three truth-values, we have to know what happens when we use logical operators in a boolean condition. The truth tables below specify the truth-value for each operator in each of the possible conditions.

A	B	A and B	A or B	A	not A
T	T	T	T	T	F
TF or FT		F	T	F	T
F	F	F	F	U	U
TU or UT		U	T		
FU or UF		F	U		
U	U	U	U		

The values in the "and", "or", and "not" columns above are common sense if you think about it.

Here's another way to look at three-valued logic. If we treat true as 1, false as 0, and unknown as 0.5, then taking the minimum of two operands gives us "and", and taking the maximum gives us "or":

A	B	as nums	A and B	min	A or B	max
T	T	1, 1	T	1	T	1
TF or FT		1, 0	F	0	T	1
F	F	0, 0	F	0	F	0
TU or UT		1, 0.5	U	0.5	T	1
FU or UF		0, 0.5	F	0	U	0.5
U	U	0.5, 0.5	U	0.5	U	0.5

Taking the complement of an operand, gives us "not":

A	as a num, x	not A	1 - x
T	1	F	0
F	0	T	1
U	0.5	U	0.5

Surprises

Recall our **Course** table:

```
csc343h-prof=> SELECT *
csc343h-prof=> FROM Course;
cnum | name | dept | breadth
```

```

cnum | name | dept | breadth
-----+-----+-----+-----
343 | Intro to Databases | CSC | f
207 | Software Design | CSC | f
148 | Intro to Comp Sci | CSC | f
263 | Data Struct & Anal | CSC | f
320 | Intro to Visual Computing | CSC | f
200 | Intro Archaeology | ANT | t
203 | Human Biol & Evol | ANT | f
150 | Organisms in Environ | EEB | f
216 | Marine Mammal Bio | EEB | f
263 | Compar Vert Anatomy | EEB | f
110 | Narrative | ENG | t
205 | Rhetoric | ENG | t
235 | The Graphic Novel | ENG | t
200 | Environmental Change | ENV | f
320 | Natl & Intl Env Policy | ENV | f
220 | Mediaeval Society | HIS | t
296 | Black Freedom | HIS | t
222 | COBOL programming | CSC | f
100 | CSC for future prime ministers | CSC | 
(19 rows)

```

Here, we ask for the courses that either have `breadth` true or have `breadth` false:

```

csc343h-prof=> SELECT *
csc343h-prof-> FROM Course
csc343h-prof-> WHERE breadth OR NOT breadth;
cnum | name | dept | breadth
-----+-----+-----+-----
343 | Intro to Databases | CSC | f
207 | Software Design | CSC | f
148 | Intro to Comp Sci | CSC | f
263 | Data Struct & Anal | CSC | f
320 | Intro to Visual Computing | CSC | f
200 | Intro Archaeology | ANT | t
203 | Human Biol & Evol | ANT | f
150 | Organisms in Environ | EEB | f
216 | Marine Mammal Bio | EEB | f
263 | Compar Vert Anatomy | EEB | f
110 | Narrative | ENG | t
205 | Rhetoric | ENG | t
235 | The Graphic Novel | ENG | t
200 | Environmental Change | ENV | f
320 | Natl & Intl Env Policy | ENV | f
220 | Mediaeval Society | HIS | t
296 | Black Freedom | HIS | t
222 | COBOL programming | CSC | f
(18 rows)

```

The expression `breadth OR NOT breadth` is a tautology, meaning it should always be true, yet we can see that **this tautology isn't true** for the course "CSC for future prime ministers" -- this course does not appear in the result!

This is a bit of a surprise, but it makes sense when we remember the truth table for `OR`:

- The truth-value of `breadth` is unknown.
- Therefore the truth-value of `NOT breadth` is unknown.
- The `OR` of two unknowns is unknown.

Here's another tautology: `AGE >= 10 OR AGE < 10`. But it's not true if `AGE` is `NULL`. Here's a little table that we can use to demonstrate this:

```
csc343h-prof=> SELECT *
csc343h-prof-> FROM Ages;
  name | age
-----+-----
  Amna | 21
  Zach | 25
  Miriam |
  Ben  | 0
(4 rows)
```

```
csc343h-prof=> SELECT *
csc343h-prof-> FROM Ages
csc343h-prof-> WHERE age >= 10 OR AGE < 10;
  name | age
-----+-----
  Amna | 21
  Zach | 25
  Ben  | 0
(3 rows)
```

Miriam is not included in the result because her age is not known. Even though `AGE >= 10 OR AGE < 10` would evaluate to true no matter what value we inserted for her age, the expression evaluates to unknown right now.

Here's one final tautology, and this time it doesn't involve any logical operators: `AGE * 0 = 0`. This mathematical rule would be true for any value of `AGE`, but it is not true in SQL when `AGE` is `NULL`, and so Miriam is again excluded:

```
csc343h-prof=> SELECT *
csc343h-prof-> FROM Ages
csc343h-prof-> WHERE AGE * 0 = 0;
  name | age
-----+-----
  Amna | 21
```

```
Zach | 25
Ben  | 0
(3 rows)
```

Summary

There are many corner conditions involving `NULL`, but you need only remember these core facts:

- Any comparison with `NULL` yields the truth-value unknown
- `WHERE` is picky: it only accepts `TRUE`.
- Therefore, `NATURAL JOIN` is picky also.
- Aggregation ignores `NULL` values.
- In other situations where an attribute may be `NULL` or where it matters whether or not two `NULL` values are considered equal, don't assume what the behaviour will be: check the documentation for your particular DBMS and test your code! And be aware that your code may not be portable to another DBMS.